

```
In [1]: import numpy as np
import pandas as pd
import os
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.metrics import calinski_harabaz_score
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)
```

```
In [3]: os.chdir("/Users/serrauzun/Desktop/MSDS_422_Practical")
df = pd.read_csv('HMEQ_Loss_clean.csv')
```

We will remove the variables with binary data as for our kmeans clustering model we should only have variables with continous data

```
In [4]: drop = df[['REASON_DebtCon', 'REASON_HomeImp', 'REASON_Missing',
'JOB_Mgr', 'JOB_Missing', 'JOB_Office', 'JOB_Other', 'JOB_ProfExe',
'JOB_Sales', 'JOB_Self']]
df = df.drop(drop, axis=1)
TARGET = df[['TARGET_BAD_FLAG']]
```

```
In [5]: x = df.copy()
x = x.drop(TARGET, axis=1)
varNames = x.columns
```

Now we will standardize our dataset with only continous data to prepare it for the model

```
In [6]: theScaler = StandardScaler()
theScaler.fit(x)
```

```
Out[6]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
In [7]: x_trn = theScaler.transform(x)
x_trn = pd.DataFrame(x_trn)
print(x_trn.head().T)
```

	0	1	2	3	4
0	-0.279351	-0.212870	-0.261452	-0.167981	-0.370407
1	-1.573958	-1.555650	-1.537343	-1.537343	-1.519036
2	-1.230735	-0.005036	-1.573541	-0.144654	0.764531
3	-1.266822	-0.640043	-1.743174	-0.195473	0.290257
4	0.244005	-0.241087	-0.656880	-0.241087	-0.795477
5	-0.281418	-0.281418	-0.281418	-0.281418	-0.281418
6	-0.373319	1.475452	-0.373319	-0.373319	-0.373319
7	-1.009413	-0.681092	-0.350779	-0.063897	-1.021765
8	-0.095518	-0.701393	-0.095518	-0.095518	-0.701393
9	-1.234969	-0.723366	-1.132648	-0.109443	-0.723366
10	-1.686771	-1.686771	-1.686771	-1.686771	-1.686771

```
In [8]: varNames_trn = []
for i in varNames:
    newName = "trn_" + i
    varNames_trn.append(newName)
print(varNames_trn)
```

```
['trn_TARGET_LOSS_AMT', 'trn_LOAN', 'trn_MORTDUE', 'trn_VALUE', 'trn_YOJ', 'trn_DEROG', 'trn_DELINQ', 'trn_CLAGE', 'trn_NINQ', 'trn_CLNO', 'trn_DEBTINC']
```

```
In [9]: x_trn.columns = varNames_trn  
print(x_trn.head().T)
```

	0	1	2	3
4				
trn_TARGET_LOSS_AMT	-0.279351	-0.212870	-0.261452	-0.167981
0407				-0.37
trn_LOAN		-1.573958	-1.555650	-1.537343
9036			-1.537343	-1.51
trn_MORTDUE		-1.230735	-0.005036	-1.573541
4531			-0.144654	0.76
trn_VALUE		-1.266822	-0.640043	-1.743174
0257			-0.195473	0.29
trn_YOJ		0.244005	-0.241087	-0.656880
5477			-0.241087	-0.79
trn_DEROG		-0.281418	-0.281418	-0.281418
1418			-0.281418	-0.28
trn_DELINQ		-0.373319	1.475452	-0.373319
3319			-0.373319	-0.37
trn_CLAGE		-1.009413	-0.681092	-0.350779
1765			-0.063897	-1.02
trn_NINQ		-0.095518	-0.701393	-0.095518
1393			-0.095518	-0.70
trn_CLNO		-1.234969	-0.723366	-1.132648
3366			-0.109443	-0.72
trn_DEBTINC		-1.686771	-1.686771	-1.686771
6771			-1.686771	-1.68

Now that we have our dataset transformed, we can start the KMeans Clustering model



We will now create lists for our clusters (K), inertia score (I), silhouette score (S) and calinski harabaz score (C). Then following creation of these empty list we will set a lower and upper limit to the number of clusters we want the model to use, in this case min 3 and max 12, and then append each score per this clustering.

```
In [10]: K_LIST = []  
I_LIST = []  
S_LIST = []  
C_LIST = []
```

```
In [14]: for K in range(3,12):
    km = KMeans(n_clusters=K, random_state = 1)
    km.fit(x_trn)
    #y = km.predict(x_trn)
    K_LIST.append(K)
    I_LIST.append(km.inertia_)
    S_LIST.append(silhouette_score(x_trn,km.labels_))
    C_LIST.append(calinski_harabaz_score(x_trn,km.labels_))
```

```
In [15]: def drawElbow(K, SCORE, LABEL):
    plt.plot(K , SCORE, 'ro-', linewidth=2)
    plt.title(LABEL)
    plt.xlabel('Clusters')
    plt.ylabel('Score')
    plt.show()
```

```
In [16]: drawElbow(K_LIST, I_LIST, "Inertia" )
drawElbow(K_LIST, S_LIST, "Silhouette" )
drawElbow(K_LIST, C_LIST, "Calinski" )
```



By looking at the Intertia, Silhouette and Calinski graphs, it seems to be the safe and reasonable approach to set the clusters from 3,12 to 4,6. The new min and max number of clusters are determined by evaluation where the elbows in the line graph occur as well as peaks and dives.

Now we write and run the clusterdata function which will give us information such as mean and count for each cluster specified

```
In [17]: def clusterData(df, x_trn, K, TARGET) :  
    print("\n\n")  
    print("K = ",K)  
    print("=====")  
    km = KMeans( n_clusters=K, random_state = 2 )  
    km.fit(x_trn)  
    Y = km.predict(x_trn)  
    df[ "CLUSTER" ] = Y  
  
    G = df.groupby( "CLUSTER" )  
    print(G.mean())  
    print("\n\n")  
    print(G[ 'TARGET_BAD_FLAG' ].value_counts())
```

As we have our function above, we can now try the number of clusters we

```
In [18]: clusterData(df, x_trn, 4, TARGET)
```

```

K = 4
=====
RTDUE \ TARGET_BAD_FLAG TARGET_LOSS_AMT LOAN MO
CLUSTER
0 0.064205 489.808186 20351.765650 54682.1
21677
1 0.943888 19510.012024 21983.366733 71695.1
25952
2 0.080000 1033.268837 23346.883721 124020.5
94102
3 0.173298 1258.192700 15060.999671 57353.7
84308

NINQ \ VALUE YOJ DEROG DELINQ CLAGE
CLUSTER
0 85388.898331 17.710594 0.081862 0.276886 243.165753
0.817014
1 101512.949900 8.257916 1.511022 2.238477 155.567166
2.352705
2 169667.082326 7.680000 0.073488 0.216744 208.958853
0.978605
3 78020.740454 5.517248 0.128576 0.220980 145.601766
1.164420

CLUSTER CLNO DEBTINC
0 22.282504 28.698192
1 26.523046 8.555556
2 26.533023 31.922441
3 17.746465 26.656674

CLUSTER TARGET_BAD_FLAG
0 0 1166
1 80
1 471
0 28
2 989
1 86
3 0 2514
1 527
Name: TARGET_BAD_FLAG, dtype: int64

```

The clustering with 4 clusters gave us the most distinct clusters. When we look at cluster 1 in our clusterData output where we used K=4, we can see that in none of the clusters the counts of 1 and 0 are too close to each other, which is what we want to see. Especially with cluster 1 we see that there are a distinct number of defaulted credits in it and when we further look into the data the mean for DEROG and DEBTINC are significantly different from the rest of the clusters. Therefore, the above clustering gives us the best results, and also the most consistent results with our other models with the same dataset.

```
In [19]: clusterData(df, x_trn, 5, TARGET)
```

K = 5					
=====					
RTDUE \ CLUSTER	TARGET_BAD_FLAG	TARGET_LOSS_AMT	LOAN	MO	
0 33020	0.084998	641.933803	16129.377203	60058.9	
1 48907	0.065817	847.424628	23966.772824	127738.1	
2 55272	0.557358	4937.345307	12381.460023	54807.4	
3 63251	0.043286	356.719081	20999.646643	55100.2	
4 58302	0.956873	22050.832884	24298.113208	76314.4	
NINQ \ CLUSTER	VALUE	YOJ	DEROG	DELINQ	CLAGE
0 1.120251	81989.342930	5.521739	0.120642	0.212691	147.026853
1 0.972399	175364.248938	7.749469	0.072187	0.203822	217.157368
2 1.457706	74675.325585	7.290556	0.222480	0.428737	149.730005
3 0.789753	86024.216484	18.010866	0.086572	0.285336	247.070751
4 2.309973	108804.799191	8.477089	1.778976	2.530997	159.556344
CLUSTER	CLNO	DEBTINC			
0	18.495104	33.759052			
1	26.728238	32.063931			
2	17.796060	0.539096			
3	22.333039	30.729035			

```
4      28.177898  10.148648
```

CLUSTER	TARGET_BAD_FLAG	
0	0	2336
	1	217
1	0	880
	1	62
2	1	481
	0	382
3	0	1083
	1	49
4	1	355
	0	16

Name: TARGET_BAD_FLAG, dtype: int64

```
In [20]: clusterData(df, x_trn, 6, TARGET)
```

K = 6	=====	TARGET_BAD_FLAG	TARGET_LOSS_AMT	LOAN	MO	
RTDUE \	CLUSTER					
0	15493	0.058887	681.556745	23946.573876	127692.7	
1	46330	0.082084	595.375296	16124.704025	59931.4	
2	19898	0.989761	24401.102389	25609.897611	81673.4	
3	75806	0.733871	11133.532258	17810.483871	60146.9	
4	01242	0.040816	319.073647	21010.825200	55073.8	
5	43581	0.558304	5003.289753	12495.759717	55426.1	
NINQ \	CLUSTER	VALUE	YOJ	DEROG	DELINQ	
0	0.966809	175363.864561	7.731263	0.074946	0.190578	217.258105
1	1.123915	81861.772691	5.514601	0.105367	0.208761	147.039339
2	2.112628	116554.626280	9.254608	0.542662	3.003413	167.051568

3	84865.649194	7.029839	4.540323	1.088710	145.439455
1.975806					
4	86024.405555	18.005590	0.082520	0.272405	247.092788
0.779059					
5	75208.223769	7.245524	0.204947	0.398115	149.643526
1.521790					

CLUSTER	CLNO	DEBTINC
0	26.669165	32.214543
1	18.492897	33.741998
2	28.160410	10.456056
3	25.193548	14.671353
4	22.309672	30.634078
5	17.904594	0.481338

CLUSTER	TARGET_BAD_FLAG	
0	0	879
	1	55
1	0	2326
	1	208
2	1	290
	0	3
3	1	91
	0	33
4	0	1081
	1	46
5	1	474
	0	375

Name: TARGET_BAD_FLAG, dtype: int64